

# Srikanth Reddy Bollampally

Agentic AI Engineer, DevOps/SRE Leader

Atlanta Metro, Open to Relocate | [srikanth.bollampally@outlook.com](mailto:srikanth.bollampally@outlook.com) | Mobile: 571-461-8282

GitHub: [github.com/mesrikanthreddy](https://github.com/mesrikanthreddy) | LinkedIn: [linkedin.com/in/bsrikanth-reddy](https://linkedin.com/in/bsrikanth-reddy) | Twitter/X: [x.com/mrbollampally](https://x.com/mrbollampally)

Hands-on engineer partnering directly with customers to deliver fast, scalable solutions across AI, infrastructure, and production deployment.

10+ Years of Experience | Azure | AWS | GCP | Generative AI | LLM Pipelines | Kubernetes | CI/CD | Agentic AI | Forward-Deployed | Multi Cloud | Fullstack | Agentic Coding

## PROFESSIONAL SUMMARY

---

Senior SRE and DevOps engineer with over a decade of experience keeping production systems reliable across GCP, Azure, and AWS. I design, operate, and tune Google Cloud databases (Cloud Spanner, Cloud SQL, and BigQuery), secure them with least-privilege IAM, and hold them to clear SLIs, SLOs, and error budgets. I lead incident response, use load testing to find and fix performance bottlenecks, and partner with engineering teams on database architecture and zero-downtime migrations. I automate the operational work with Terraform and scripting, and bring the same customer-focused, fast-moving execution to every team I support.

## CORE SKILLS

---

- **Core Skills:** Focus on foundational professional abilities and experience.
- **AI, DevOps & CloudOps:** Leverage cloud technologies and modern development practices.
- **Customer Focus:** Prioritize customer needs with a dedicated, service-oriented approach.

## KEY HIGHLIGHTS

---

- **10 Years in Customer-Facing Delivery:** Cloud, DevOps, and AI Engineering across enterprise environments with direct customer collaboration
- **Faster Inference:** Reduced latency through rapid model optimization, caching, and batching to improve customer experience
- **Cost Savings:** Lowered cloud compute spend through fast iteration on infrastructure optimization and right-sizing
- **Engineers Unblocked:** Mentored across 2 teams with reusable AI patterns, reference architectures, and hands-on support

## PROFESSIONAL EXPERIENCE

---

### YTT Global • Stealth Healthcare

*Principal AI Engineer / Principal Technical Consultant* | Jan 2026 – Present

- Partnered directly with stakeholders to take cloud, data, and AI/ML initiatives from ambiguous problem statements to production-ready solutions.
- Worked across engineering, security/risk, and business teams to clarify requirements, guide architecture decisions on GCP data services, IAM, and platform reliability, and define practical delivery plans that kept execution moving.
- Led rapid problem-solving in customer-facing, high-stakes environments by translating complex technical tradeoffs into clear options, executive-ready narratives, and release-readiness criteria.
- Focused on tangible business outcomes, dependable production delivery, and systems that stay reliable at scale.

### Key Deliverables

- Multi-agent LLM systems using MCP and tool-calling orchestration with deterministic-validation agentic framework (confidence thresholds + Human-in-the-Loop gates) for reliable, auditable autonomous execution
- RAG pipelines, semantic chunking, embedding, retrieval, reranking, and grounded generation, using Databricks and pgvector, with monitoring for retrieval quality, groundedness, hallucination rate, and drift
- VLM-powered document processing cutting manual review effort 50%

- Responsible AI controls PII tokenization proxy (GDPR-aligned), prompt-injection defenses, authorization boundaries for secure tool execution
- Classical machine learning (ML) and artificial intelligence (AI) solutions encompass a wide range of applications, including classification, regression, clustering, natural language processing (NLP), and computer vision. These solutions leverage statistical modeling, feature engineering, and algorithm selection to enhance predictive accuracy, automate decision-making processes, and support production use cases beyond the capabilities of generative artificial intelligence (GenAI).

### **Georgia Technology Authority**

*Lead DevOps Engineer, MLOps | Feb 2025 – Dec 2025*

- Worked directly with state stakeholders to turn evolving priorities into shippable delivery plans, translating ambiguous requirements into practical migration and automation decisions.
- Led the state's app portfolio migration from AWS CodeCommit to Bitbucket, and built Terraform-based deployment workflows for ECS, EKS, and EC2.
- Migrated SharePoint workloads from Azure to AWS, with the infrastructure provisioned as code in Terraform.
- Wrote technical documentation, including runbooks and architecture guides, so teams could operate and extend the new platforms.
- Partnered with engineering and the Georgia Innovation Lab to build AI/ML proofs of concept, including agentic AI for policy-document summarization and other GenAI use cases, focused on tangible operational impact.

### **K-Force | Oak Street Health**

*Lead DevOps/SRE, ML Ops | Oct 2023 – Feb 2025*

- Served as SRE for patient-facing .NET applications on Azure, with Cisco Meraki networking, in a high-stakes healthcare environment.
- Defined SLIs, SLOs, and error budgets with application teams to consistently meet SLA commitments, and built Grafana dashboards that, together with ThousandEyes, gave teams clear visibility into service health and network performance.
- Ran k6 and JMeter load and performance tests to pinpoint bottlenecks, then partnered with developers to resolve them and reduce application slowness.
- Supported application migrations by giving developers SRE feedback on reliability and performance during development, automated integration workflows with Workato, and delivered zero-downtime Kubernetes upgrades.
- fixed the application p95 response time going from 2s seconds to 400ms.

### **Infosys | Cisco & UnitedHealth Group**

*Lead DevOps & Senior Cloud Engineer | 2019 – 2023*

- Embedded with customer teams to deliver cloud and data platform outcomes across complex enterprise environments.
- Designed and managed Google Cloud Spanner and Cloud SQL databases for production workloads, covering schema design, query tuning, and performance monitoring, and enforced least-privilege access with GCP IAM roles and service accounts.
- Partnered across engineering, operations, and business stakeholders to build data pipelines on Snowflake, BigQuery, Prefect, and Kubernetes that improved efficiency 45%, and shipped pipeline-as-code CI/CD automation that reduced operational costs 30%.
- Solved large-scale migration and stability challenges quickly, optimizing 1,000+ EC2 instances and achieving 99.99% system stability with zero-downtime cloud migration.

### **Steel Thread | USPTO**

*DevOps Engineer | 2016 – 2019*

- Worked side by side with stakeholders to modernize legacy systems in a government environment with shifting requirements and tight delivery constraints.
- Built a Docker proof of concept that validated containerization for legacy applications, then led the migration from Mesos to on-prem Kubernetes, improving scalability 40% and cutting deployment complexity 50%.

- Migrated CI/CD from open-source Jenkins to CloudBees, and automated server configuration and environment provisioning with Ansible.
- Shipped DevSecOps automation and TLS/SSL hygiene improvements that reduced environment setup time 50% and lowered security vulnerabilities 50%, delivering clearer, faster, and more secure delivery for the organization.

## FORWARD-DEPLOYED STRENGTHS

---

- **Customer Communication:** Translates technical complexity into clear options, tradeoffs, and next steps for customers and stakeholders
- **Rapid Prototyping:** Builds working proofs of concept fast to validate ideas, unblock decisions, and shorten feedback loops
- **Cross-Functional Collaboration:** Partners with engineering, product, security, and operations teams to move from ambiguity to shipped outcomes
- **Problem Solving Under Ambiguity:** Owns messy, high-context problems and breaks them into practical actions that deliver value quickly
- **Shipping Fast:** Uses tight iteration cycles, strong execution, and hands-on delivery to turn customer needs into production impact

## TECHNICAL SKILLS

---

- **DevOps & CI/CD:** GitHub Actions, GitLab CI, Jenkins, ArgoCD, Terraform, Ansible, Pulumi
- **Cloud Engineering:** AWS, GCP, Azure — VPC, IAM, Serverless, Kubernetes (AKS/EKS/GKE)
- **AI & Machine Learning:** Python, PyTorch, TensorFlow, LLMs, RAG, Fine-Tuning, Transformers, NLP, Computer Vision
- **Observability:** Prometheus, Grafana, Datadog, OpenTelemetry, ELK Stack, Chaos Engineering, SLO/Error Budget Management

## CERTIFICATIONS

---

- Terraform Certified Associate — Feb 2024
- LFS162: Introduction to DevOps & SRE — Aug 2024
- LFS116: PyTorch & Deep Learning for Decision Makers — Sep 2024
- Foundations of Project Management — Aug 2023
- Intensive Program — Machine Learning & AI, Deep Atlas — Aug 2024

## PERSONAL PROJECTS

---

- **kshield** — Local-first AI code review firewall catching secrets, broken access control, and AI hallucinations on-device
- **qga-cancer-therapeutics** — Quantum Genetic Algorithm POC for multi-drug cancer therapy optimization (NVIDIA CUDA-Q)

## VOLUNTEERING

---

- Habitat for Humanity International (2017–Present)
- American Red Cross Disaster & Human Relief (2026–Present)

## LANGUAGES

---

English (Fluent) · Telugu (Fluent) · Hindi (Conversational) · Spanish (Basic)